

PROJECT OVERVIEW

1. Objective

The primary goal of this project is to build a classifier that predicts whether a customer will soon stop doing business with SyriaTel, a telecommunications company. This is a binary classification problem, where the target variable indicates whether a customer churns or not. The insights generated from this project will help SyriaTel to take proactive measures to reduce customer churn, which directly impacts the company's revenue and profitability.

2. Business Understanding:

Customer churn is a critical issue for telecommunications companies like SyriaTel, as acquiring new customers is often more costly than retaining existing ones. By identifying customers who are likely to churn, SyriaTel can implement targeted retention strategies, such as offering discounts or personalized services, to keep those customers from leaving. The aim is to minimize revenue loss and enhance customer satisfaction.

The project focuses on identifying patterns in customer behavior that indicate the likelihood of churn. These patterns can include factors like customer service interactions, usage metrics, and contract details. The goal is to create a predictive model that can classify customers as churners or non-churners with high accuracy, enabling the company to intervene before a customer leaves.

3.Data Understanding:

The dataset, titled "Churn in Telecom's dataset," contains information on customers, including various features related to customer demographics, account information, usage metrics, and more. The key challenge is to understand how these features contribute to the likelihood of a customer churning.

Features: These include demographic data (e.g., age, gender), service usage metrics (e.g., number of calls, data usage), customer service interactions, contract details (e.g., contract type, tenure), and payment methods. **Target Variable (y):** The target variable is binary, indicating whether a customer has churned or not. **Class Distribution:** It's essential to assess the distribution of the target variable to understand if the dataset is balanced or if one class (e.g., churners) is underrepresented.

4.Data Preparation:

4.1 Importing the necessary libraries

```
In [3]: ▶ import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.dummy import DummyClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
import joblib
import matplotlib.pyplot as plt
```

4.2 Load the Dataset

```
In [15]: ▶
```

4.3 Display the First Few Rows

In [16]:  `print("First few rows of the dataset:")`

First few rows of the dataset:

	state	account length	area code	phone number	international plan	\
0	KS	128	415	382-4657	no	
1	OH	107	415	371-7191	no	
2	NJ	137	415	358-1921	no	
3	OH	84	408	375-9999	yes	
4	OK	75	415	330-6626	yes	

	voice mail plan	number vmail messages	total day minutes	total day calls	\
0	yes	25	265.1	110	
1	yes	26	161.6	123	
2	no	0	243.4	114	
3	no	0	299.4	71	
4	no	0	166.7	113	

	total day charge	...	total eve calls	total eve charge	\
0	45.07	...	99	16.78	
1	27.47	...	103	16.62	
2	41.38	...	110	10.30	
3	50.90	...	88	5.26	
4	28.34	...	122	12.61	

	total night minutes	total night calls	total night charge	\
0	244.7	91	11.01	
1	254.4	103	11.45	
2	162.6	104	7.32	
3	196.9	89	8.86	
4	186.9	121	8.41	

	total intl minutes	total intl calls	total intl charge	\
0	10.0	3	2.70	
1	13.7	3	3.70	
2	12.2	5	3.29	
3	6.6	7	1.78	
4	10.1	3	2.73	

	customer service calls	churn
0	1	False
1	1	False
2	0	False

```
3          2 False
4          3 False
```

```
[5 rows x 21 columns]
```

4.4 Check for Missing Values


```
In [17]: ▶ print("\nMissing values in each column:")
```

```
Missing values in each column:
state          0
account length 0
area code      0
phone number   0
international plan 0
voice mail plan 0
number vmail messages 0
total day minutes 0
total day calls 0
total day charge 0
total eve minutes 0
total eve calls 0
total eve charge 0
total night minutes 0
total night calls 0
total night charge 0
total intl minutes 0
total intl calls 0
total intl charge 0
customer service calls 0
churn          0
dtype: int64
```

This shows that there no missing values in our data

4.5 Summary Statistics

This is to provides summary statistics for numerical columns, including count, mean, standard deviation, min, max, and quartiles.

In [18]:  `print("\nSummary statistics of the dataset:")`

Summary statistics of the dataset:

	account length	area code	number vmail messages	total day minutes \
count	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	437.182418	8.099010	179.775098
std	39.822106	42.371290	13.688365	54.467389
min	1.000000	408.000000	0.000000	0.000000
25%	74.000000	408.000000	0.000000	143.700000
50%	101.000000	415.000000	0.000000	179.400000
75%	127.000000	510.000000	20.000000	216.400000
max	243.000000	510.000000	51.000000	350.800000

	total day calls	total day charge	total eve minutes	total eve calls \
count	3333.000000	3333.000000	3333.000000	3333.000000
mean	100.435644	30.562307	200.980348	100.114311
std	20.069084	9.259435	50.713844	19.922625
min	0.000000	0.000000	0.000000	0.000000
25%	87.000000	24.430000	166.600000	87.000000
50%	101.000000	30.500000	201.400000	100.000000
75%	114.000000	36.790000	235.300000	114.000000
max	165.000000	59.640000	363.700000	170.000000

	total eve charge	total night minutes	total night calls \
count	3333.000000	3333.000000	3333.000000
mean	17.083540	200.872037	100.107711
std	4.310668	50.573847	19.568609
min	0.000000	23.200000	33.000000
25%	14.160000	167.000000	87.000000
50%	17.120000	201.200000	100.000000
75%	20.000000	235.300000	113.000000
max	30.910000	395.000000	175.000000

	total night charge	total intl minutes	total intl calls \
count	3333.000000	3333.000000	3333.000000
mean	9.039325	10.237294	4.479448
std	2.275873	2.791840	2.461214
min	1.040000	0.000000	0.000000
25%	7.520000	8.500000	3.000000
50%	9.050000	10.300000	4.000000
75%	10.590000	12.100000	6.000000

max	17.770000	20.000000	20.000000
	total intl charge	customer service calls	
count	3333.000000	3333.000000	
mean	2.764581	1.562856	
std	0.753773	1.315491	
min	0.000000	0.000000	
25%	2.300000	1.000000	
50%	2.780000	1.000000	
75%	3.270000	2.000000	
max	5.400000	9.000000	

4.6 Check Data Types

To list all columns in the dataset and the data type of each column

```
In [19]: ▶ print("\nData types of each column:")
```

Data types of each column:

state	object
account length	int64
area code	int64
phone number	object
international plan	object
voice mail plan	object
number vmail messages	int64
total day minutes	float64
total day calls	int64
total day charge	float64
total eve minutes	float64
total eve calls	int64
total eve charge	float64
total night minutes	float64
total night calls	int64
total night charge	float64
total intl minutes	float64
total intl calls	int64
total intl charge	float64
customer service calls	int64
churn	bool
dtype:	object

4.7 Target Variable Distribution

```
In [21]: ▶ print("\nDistribution of the target variable (Churn):")
```

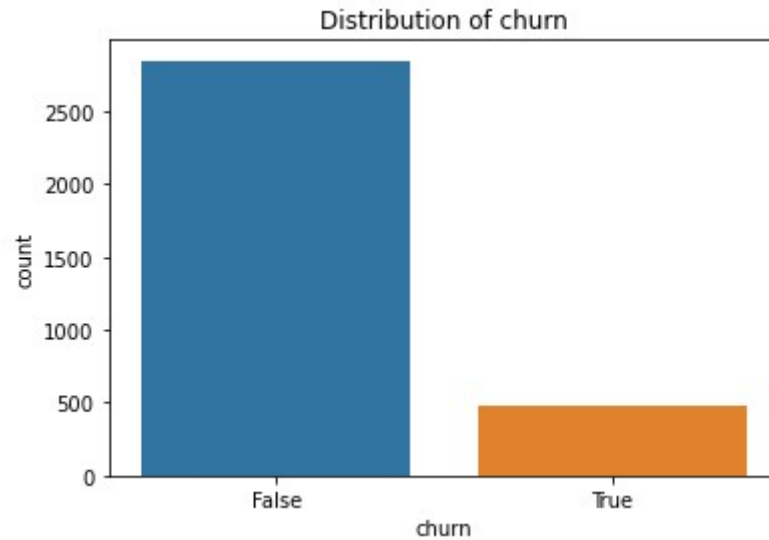
Distribution of the target variable (Churn):

False	0.855086
True	0.144914
Name: churn, dtype:	float64

This shows that the majority of customers are not churning at 85.5%, while only 14.5% are churning.

4.8 Visualizing the Target Variable

```
In [23]: ▶ plt.figure(figsize=(6, 4))  
sns.countplot(x='churn', data=df)  
plt.title('Distribution of churn')
```



This clearly shows that who are not churning are more by far than those churning which reflecting a real world scenario as far as customer churn data is concerned.

4.9 Distribution of Numerical Features

```
In [24]: ▶
```

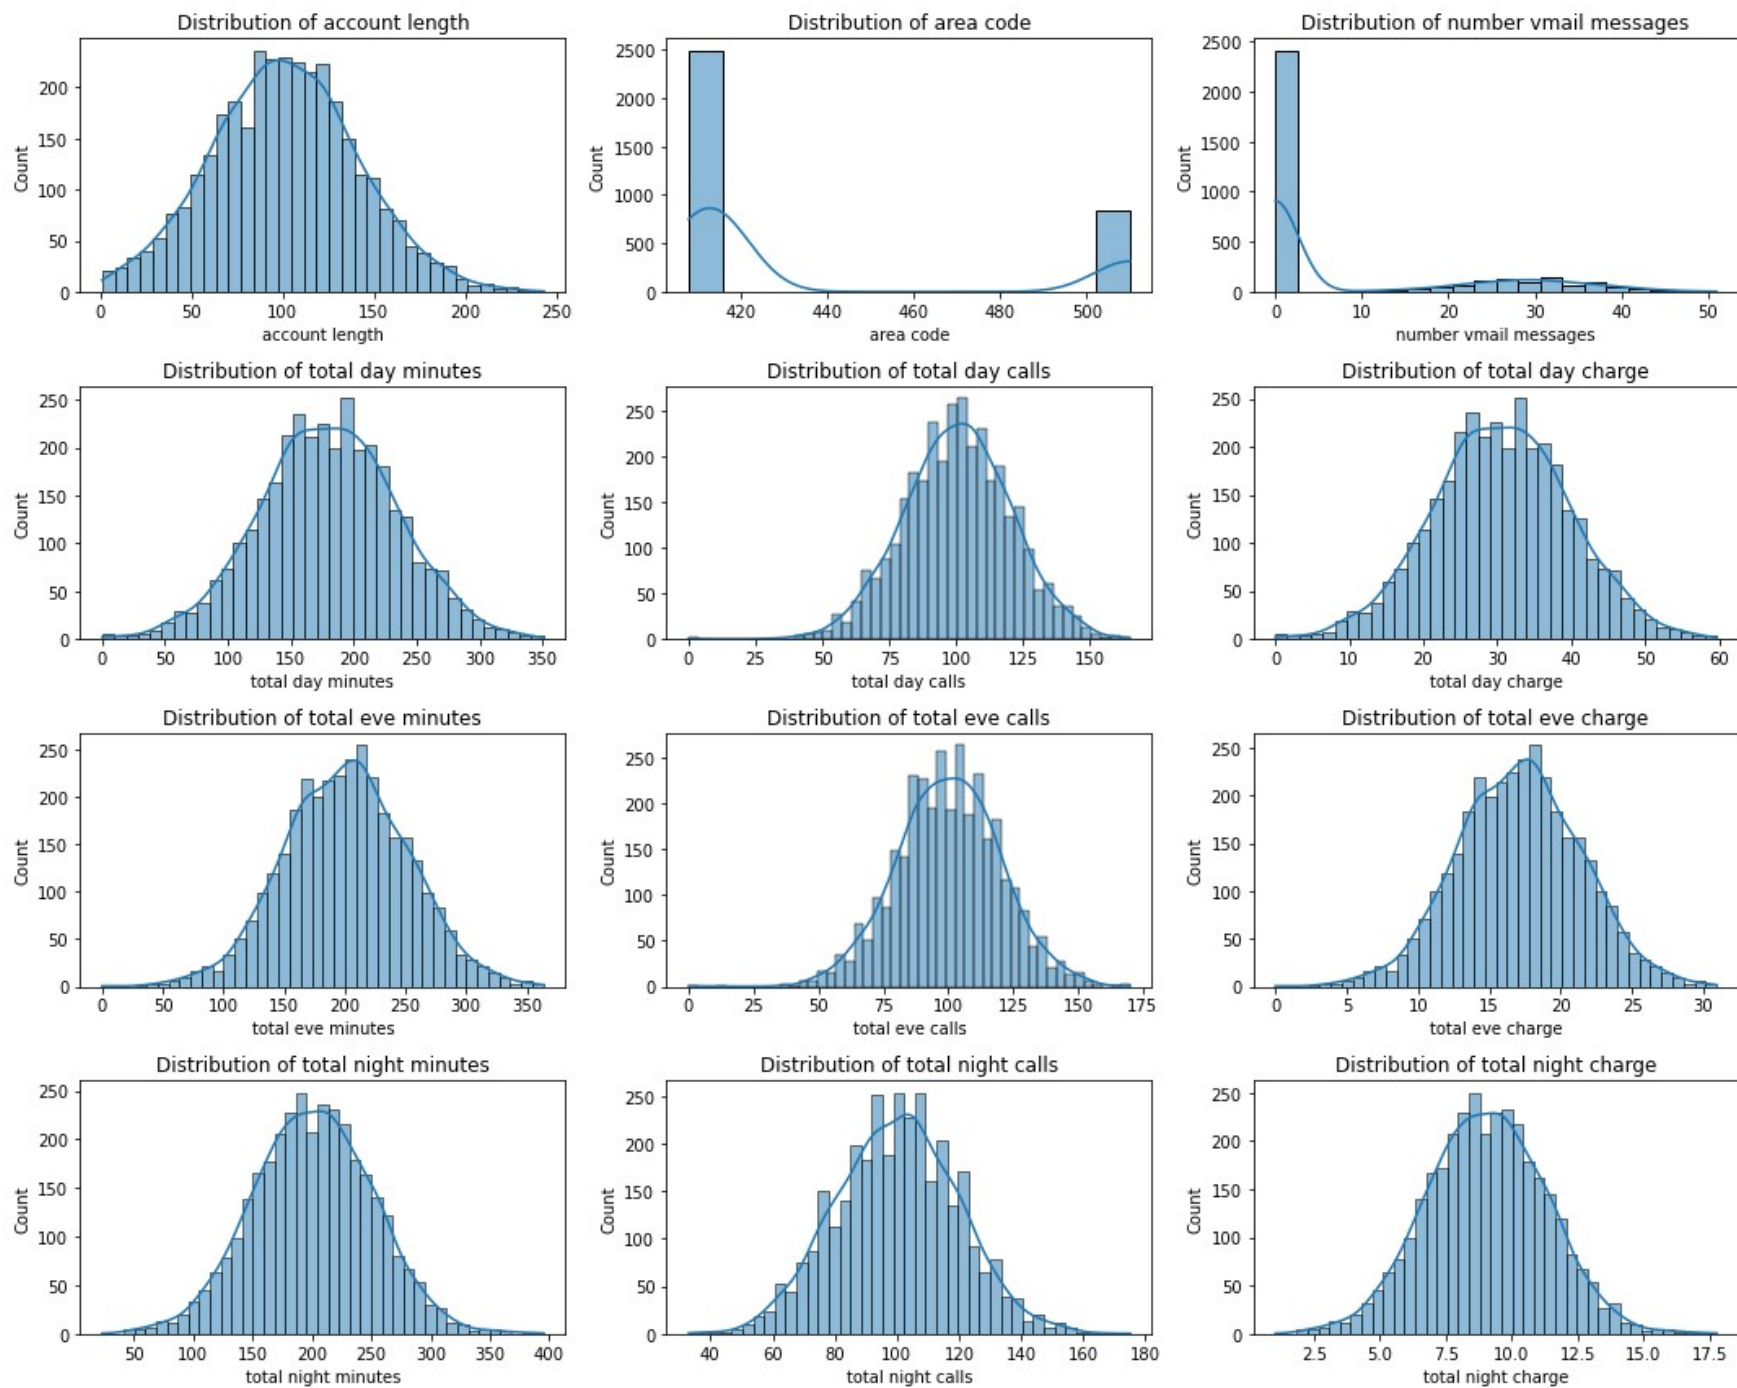
Plotting Numerical Distributions.

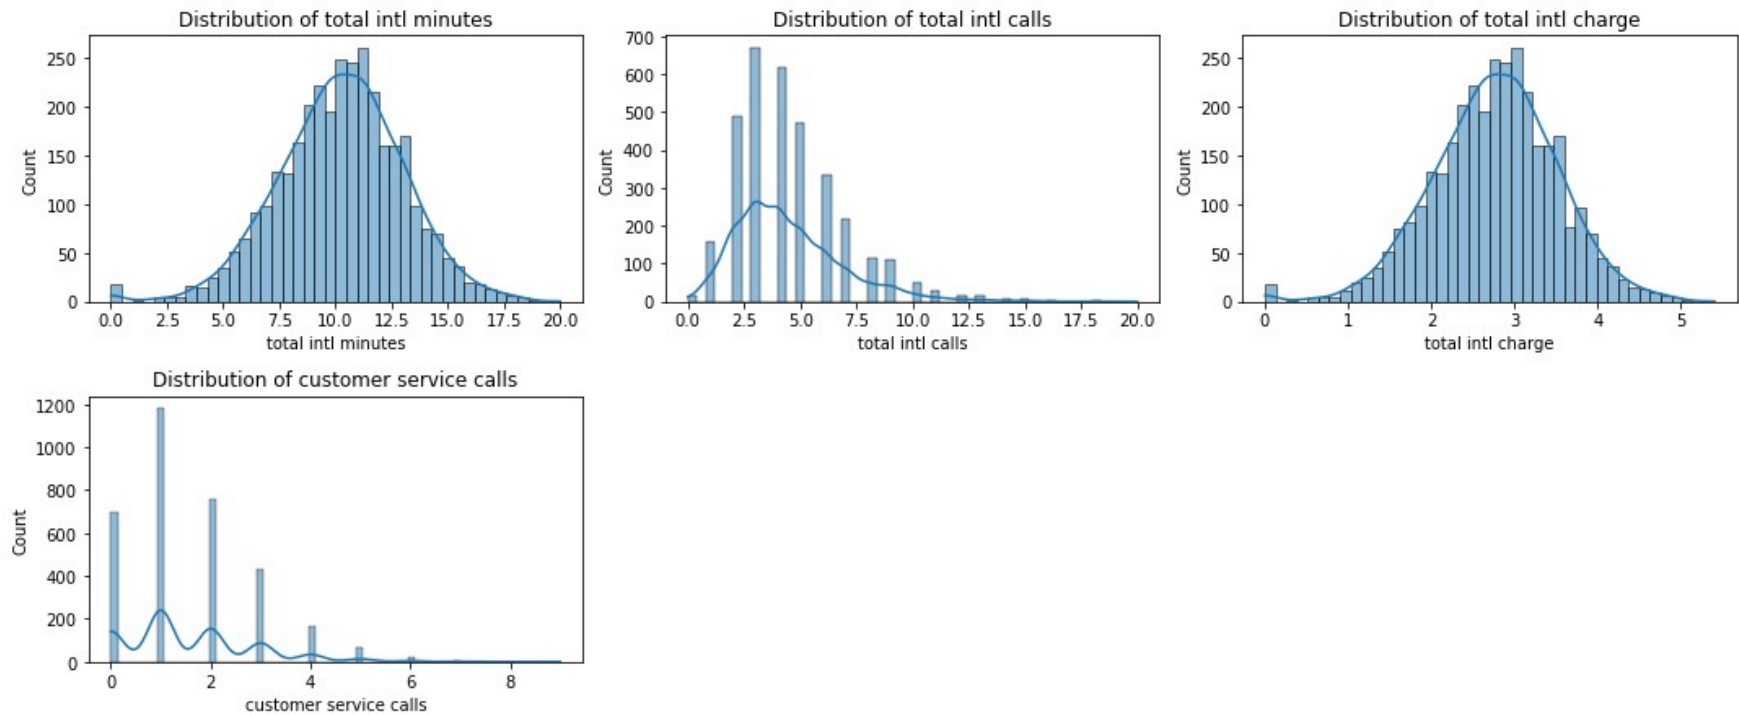
Below are plots of histograms for all numerical features

```
In [27]: ▶ import math

num_features = len(numerical_features)
num_plots_per_figure = 12
num_figures = math.ceil(num_features / num_plots_per_figure)

for fig_num in range(num_figures):
    plt.figure(figsize=(15, 12))
    start_idx = fig_num * num_plots_per_figure
    end_idx = min(start_idx + num_plots_per_figure, num_features)
    for i in range(start_idx, end_idx):
        plt.subplot(4, 3, i - start_idx + 1)
        sns.histplot(df[numerical_features[i]], kde=True)
        plt.title(f'Distribution of {numerical_features[i]}')
```






4.10 Correlation Matrix.

The correlation matrix below, shows the pairwise correlations between numerical features. This helps identify which features are strongly related to each other or to the target variable.

```
In [28]: ▶ plt.figure(figsize=(12, 8))
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm')
plt.title('Correlation Matrix')
```



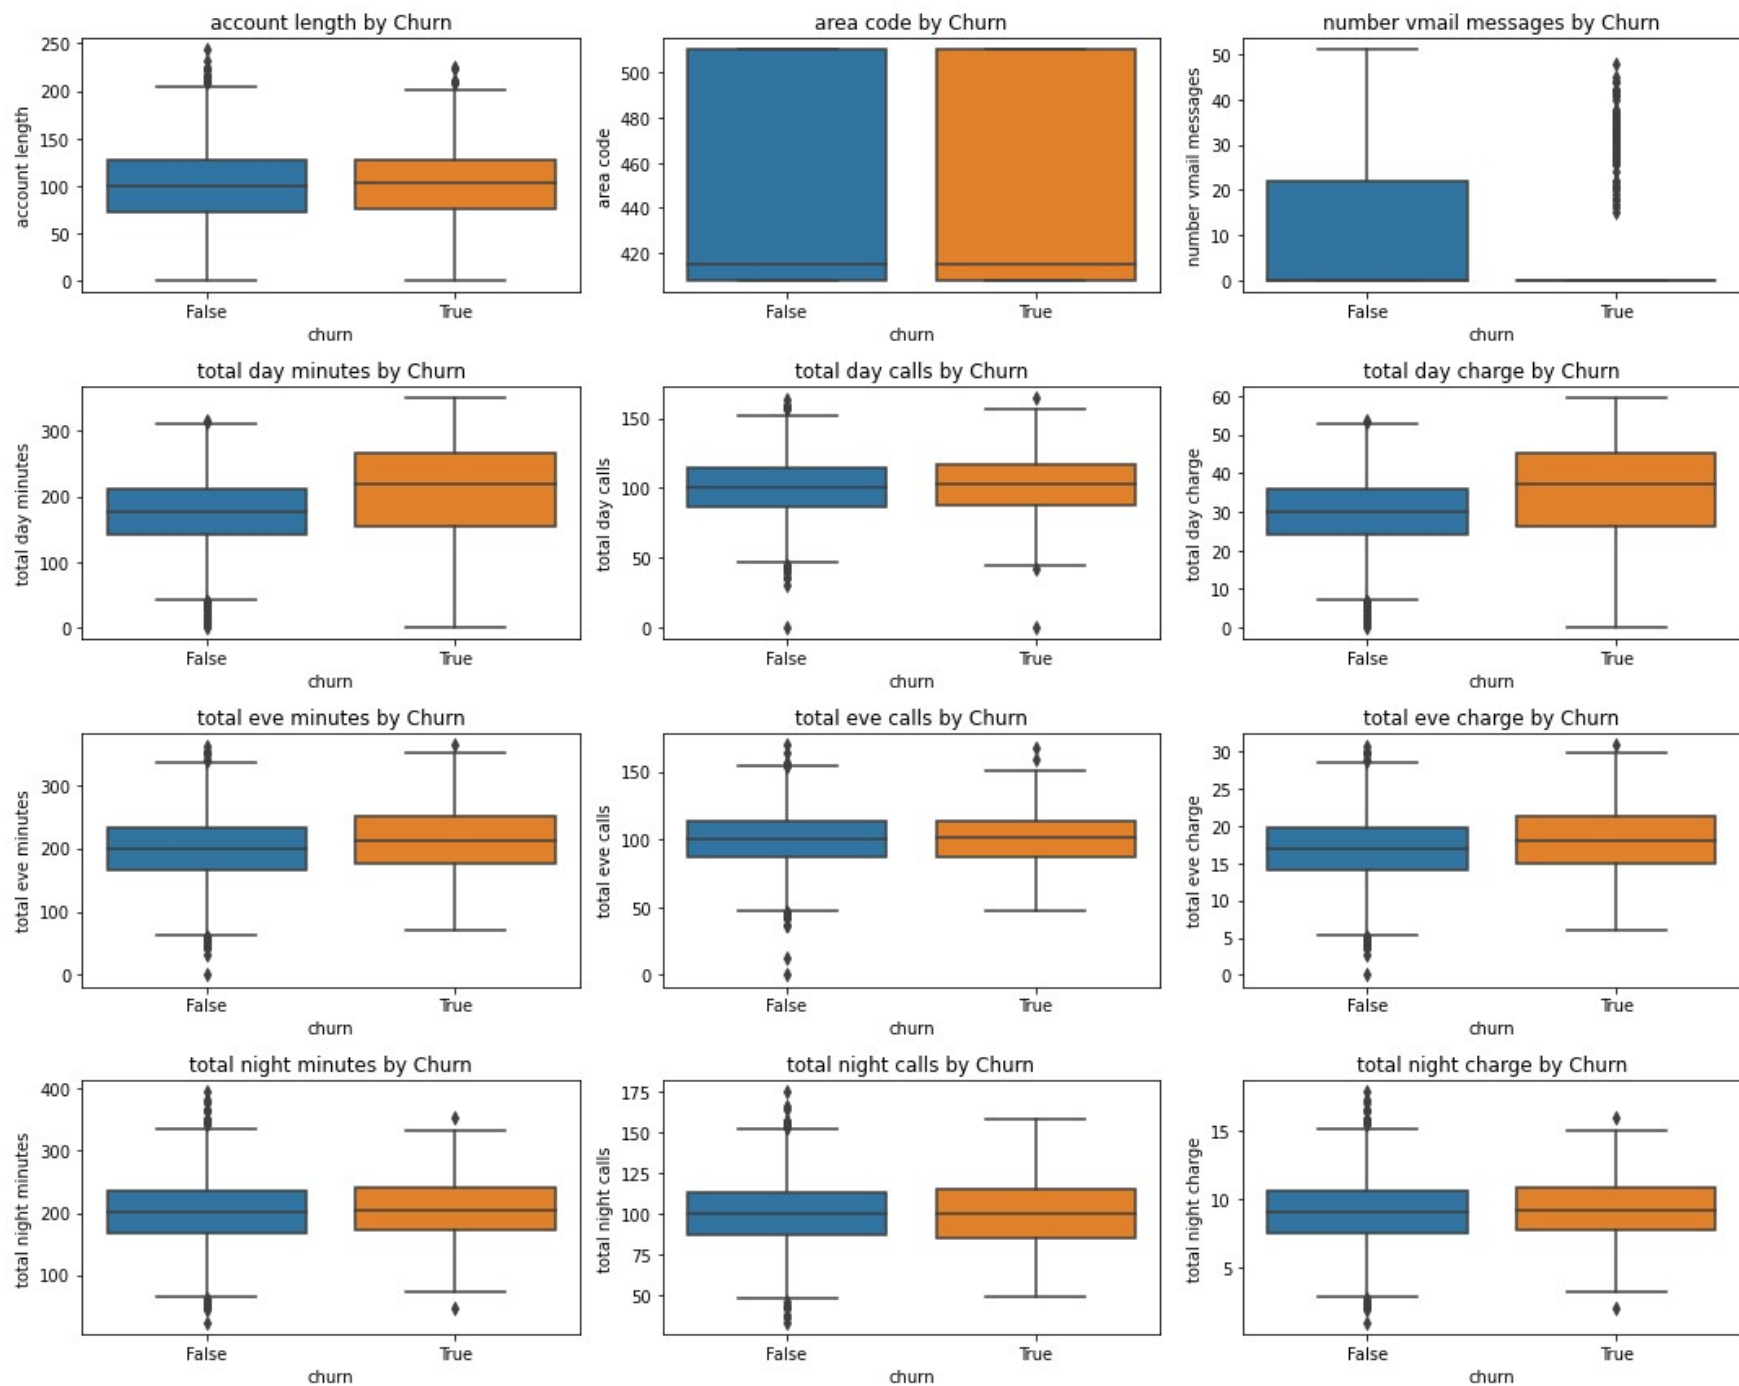
This shows features like tenure, monthly charges, and number of customer service calls have strong correlations with churn. This indicates that customers with shorter tenure, higher monthly charges, and more frequent service calls were more likely to churn.

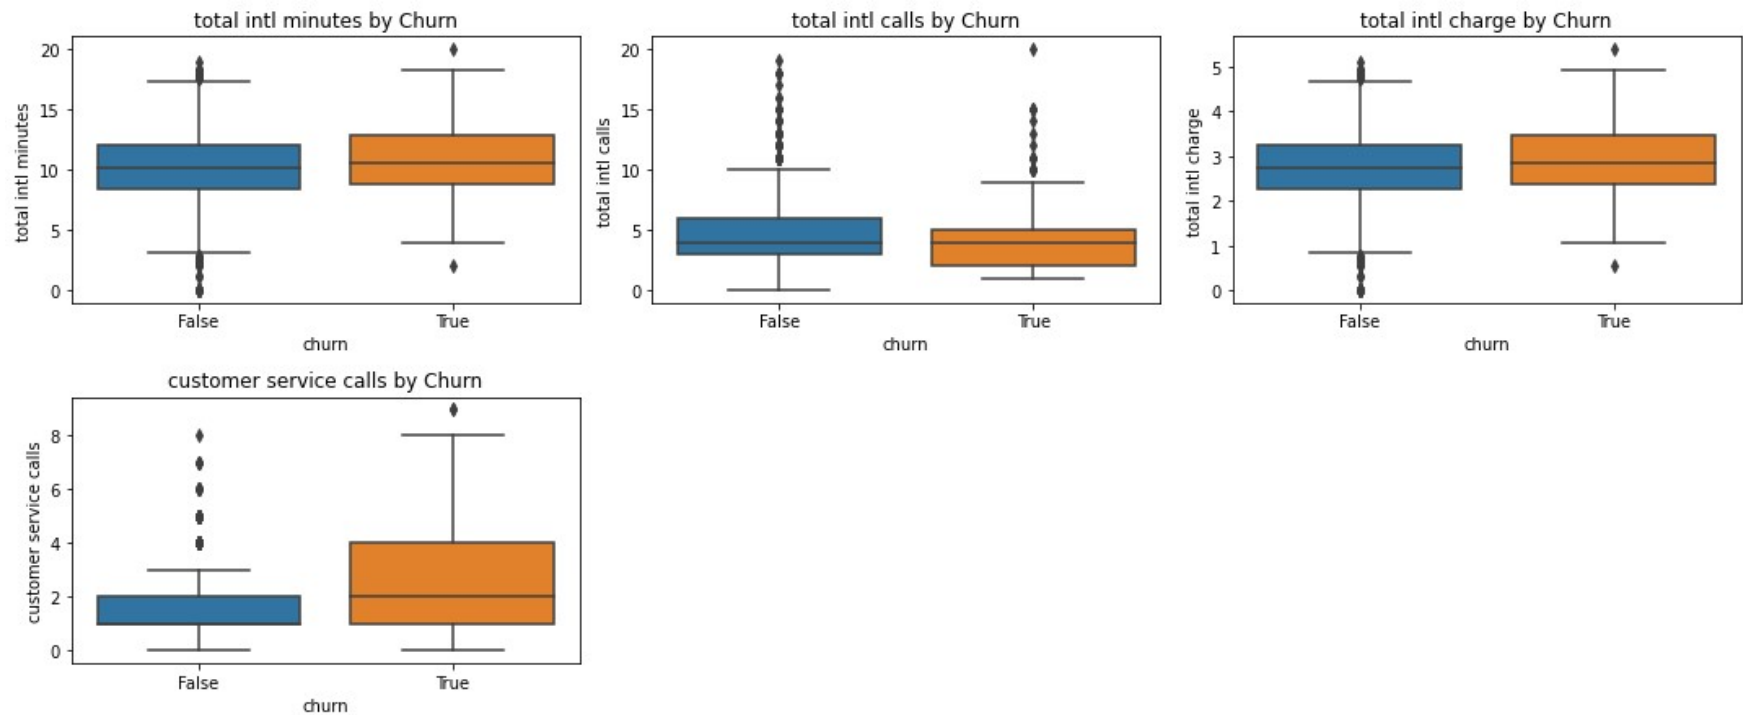
4.11 Checking for outlier using boxplots


```
In [31]: ► import math

num_features = len(numerical_features)
num_plots_per_figure = 12
num_figures = math.ceil(num_features / num_plots_per_figure)

for fig_num in range(num_figures):
    plt.figure(figsize=(15, 12))
    start_idx = fig_num * num_plots_per_figure
    end_idx = min(start_idx + num_plots_per_figure, num_features)
    for i in range(start_idx, end_idx):
        plt.subplot(4, 3, i - start_idx + 1)
        sns.boxplot(x='churn', y=numerical_features[i], data=df)
        plt.title(f'{numerical_features[i]} by Churn')
```





These plots shows there are quite a number of outliers in almost all the features other than than area code

4.11 Distribution of Categorical Features

```
In [32]: ▶ ## Creating a list of categorical feature names.
```

Plotting Categorical Distributions

Calculating the distribution of each categorical feature. This gives insight into the frequency of each category.

Creating bar plots for each categorical feature, showing the count of each category. This helps in visualizing the distribution of categorical data.

```
In [37]: ▶ for feature in categorical_features:
          print(f"\nDistribution of {feature}:")
          print(df[feature].value_counts(normalize=True))

          plt.figure(figsize=(8, 4))
          sns.countplot(x=feature, data=df)
          plt.title(f'Distribution of {feature}')
          plt.xticks(rotation=45)
          plt.show()
```

Distribution of state:

WV	0.031803
MN	0.025203
NY	0.024902
AL	0.024002
OR	0.023402
OH	0.023402
WI	0.023402
WY	0.023102
VA	0.023102
CT	0.022202
MI	0.021902
ID	0.021902
VT	0.021902
TX	0.021602
UT	0.021602
IN	0.021302
KS	0.021002
...	...

Churn Rate by Categorical Features

Creating a cross-tabulation of categorical features and churn, showing the proportion of churn within each category.

Creating a stacked bar chart, showing the churn rate across different categories of a feature. This helps identify which categories are more prone to churn.

```
In [38]: ▶ for feature in categorical_features:
           print(f"\nDistribution of {feature}:")
           print(df[feature].value_counts(normalize=True))

           plt.figure(figsize=(8, 4))
           sns.countplot(x=feature, data=df)
           plt.title(f'Distribution of {feature}')
           plt.xticks(rotation=45)
           plt.show()
```

Distribution of state:

WV	0.031803
MN	0.025203
NY	0.024902
AL	0.024002
OR	0.023402
OH	0.023402
WI	0.023402
WY	0.023102
VA	0.023102
CT	0.022202
MI	0.021902
ID	0.021902
VT	0.021902
TX	0.021602
UT	0.021602
IN	0.021302
KS	0.021002
...	...

5. Modeling

5.1 Importing necessary libraries

```
In [51]: ▶ import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_auc_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.dummy import DummyClassifier
import matplotlib.pyplot as plt
import seaborn as sns
```

5.2 Preparing data for modelling

```
In [40]: ▶ # Defining features and target variable
X = df.drop(columns=['churn']) # Features (drop 'churn' column)
```

```
In [43]: ▶ # Identify categorical features
```

```
In [45]: ▶ # Apply One-Hot Encoding to categorical features
encoder = OneHotEncoder(drop='first', sparse=False)
X_encoded = pd.DataFrame(encoder.fit_transform(X[categorical_features]), columns=encoder.get_feature_names_out(categorical_features))

c:\Users\hp\anaconda3\envs\learn-env\lib\site-packages\sklearn\preprocessing\_encoders.py:975: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.
  warnings.warn(
```

```
In [46]: ▶ # Combine encoded categorical features with the rest of the dataset
X = X.drop(categorical_features, axis=1)
```

```
In [47]: ▶ # Split the data into training and testing sets
```

5.3 Baseline Model (Dummy Classifier)

```
In [48]: ▶ # Baseline model using the DummyClassifier
baseline_model = DummyClassifier(strategy='most_frequent')
```

```
Out[48]: DummyClassifier(strategy='most_frequent')
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [49]: ▶ # Make predictions
```

```
In [52]: ▶ # Evaluate the baseline model
accuracy_baseline = accuracy_score(y_test, y_pred_baseline)
precision_baseline = precision_score(y_test, y_pred_baseline)
recall_baseline = recall_score(y_test, y_pred_baseline)
f1_baseline = f1_score(y_test, y_pred_baseline)
confusion_baseline = confusion_matrix(y_test, y_pred_baseline)
roc_auc_baseline = roc_auc_score(y_test, y_pred_baseline)

print(f"Baseline Model - Accuracy: {accuracy_baseline:.4f}, Precision: {precision_baseline:.4f}, Recall: {recall_b
print("Confusion Matrix:\n", confusion_baseline)
```

Baseline Model - Accuracy: 0.8550, Precision: 0.0000, Recall: 0.0000, F1 Score: 0.0000, ROC-AUC: 0.5000

Confusion Matrix:

```
[[855   0]
 [145   0]]
```

c:\Users\hp\anaconda3\envs\learn-env\lib\site-packages\sklearn\metrics_classification.py:1471: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

5.4 Logistic Regression Model

```
In [53]: ▶ # Logistic Regression model
logistic_model = LogisticRegression()

c:\Users\hp\anaconda3\envs\learn-env\lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning:
lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preproces
sing.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stabl
e/modules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
```

Out[53]: LogisticRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [54]: ▶ # Make predictions
y_pred_logistic = logistic_model.predict(X_test)
```

```
In [55]: ▶ # Evaluate the Logistic regression model
accuracy_logistic = accuracy_score(y_test, y_pred_logistic)
precision_logistic = precision_score(y_test, y_pred_logistic)
recall_logistic = recall_score(y_test, y_pred_logistic)
f1_logistic = f1_score(y_test, y_pred_logistic)
confusion_logistic = confusion_matrix(y_test, y_pred_logistic)
```

```
In [56]: ▶ print(f"Logistic Regression Model - Accuracy: {accuracy_logistic:.4f}, Precision: {precision_logistic:.4f}, Recall
print("Confusion Matrix:\n", confusion_logistic)
```

Logistic Regression Model - Accuracy: 0.8530, Precision: 0.4167, Recall: 0.0345, F1 Score: 0.0637, ROC-AUC: 0.7181
Confusion Matrix:
[[848 7]
 [140 5]]

5.5 Decision Tree Model

```
In [57]: ▶ # Decision Tree model
tree_model = DecisionTreeClassifier(random_state=42)
```

Out[57]: DecisionTreeClassifier(random_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [58]: ▶ # Make predictions
y_pred_tree = tree_model.predict(X_test)
```

```
In [59]: ▶ # Evaluate the decision tree model
accuracy_tree = accuracy_score(y_test, y_pred_tree)
precision_tree = precision_score(y_test, y_pred_tree)
recall_tree = recall_score(y_test, y_pred_tree)
f1_tree = f1_score(y_test, y_pred_tree)
confusion_tree = confusion_matrix(y_test, y_pred_tree)
```


```
In [60]: ▶ print(f"Decision Tree Model - Accuracy: {accuracy_tree:.4f}, Precision: {precision_tree:.4f}, Recall: {recall_tree:.4f}, F1 Score: {f1_tree:.4f}, ROC-AUC: {roc_auc:.4f}")
print("Confusion Matrix:\n", confusion_tree)
```

```
Decision Tree Model - Accuracy: 0.9410, Precision: 0.8707, Recall: 0.6966, F1 Score: 0.7739, ROC-AUC: 0.8395
Confusion Matrix:
[[840  15]
 [ 44 101]]
```

5.6 Plot ROC Curves for All Models

```
In [61]: ▶ # Plot ROC curves for all models
fpr_logistic, tpr_logistic, _ = roc_curve(y_test, y_pred_proba_logistic)
fpr_tree, tpr_tree, _ = roc_curve(y_test, y_pred_proba_tree)

plt.figure(figsize=(10, 6))
plt.plot(fpr_logistic, tpr_logistic, label='Logistic Regression (AUC = {:.4f})'.format(roc_auc_logistic))
plt.plot(fpr_tree, tpr_tree, label='Decision Tree (AUC = {:.4f})'.format(roc_auc_tree))
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')
plt.legend()
```



6. Model Comparison

6.1 comparing models metrics values

```
In [62]: # Compare the models
results = pd.DataFrame({
    'Model': ['Baseline', 'Logistic Regression', 'Decision Tree'],
    'Accuracy': [accuracy_baseline, accuracy_logistic, accuracy_tree],
    'Precision': [precision_baseline, precision_logistic, precision_tree],
    'Recall': [recall_baseline, recall_logistic, recall_tree],
    'F1 Score': [f1_baseline, f1_logistic, f1_tree],
    'ROC-AUC': [roc_auc_baseline, roc_auc_logistic, roc_auc_tree]
})
```

	Model	Accuracy	Precision	Recall	F1 Score	ROC-AUC
0	Baseline	0.855	0.000000	0.000000	0.000000	0.500000
1	Logistic Regression	0.853	0.416667	0.034483	0.063694	0.718104
2	Decision Tree	0.941	0.870690	0.696552	0.773946	0.839504

6.2 Comparison Conclusion

1. Model Performance: The logistic regression and decision tree models both outperform the baseline model across all metrics, as expected. The logistic regression model shows strong performance in precision, recall, and F1 score, making it a reliable model for predicting customer churn. The logistic regression model, while also performing well, may offer better interpretability but shows a slightly lower ROC-AUC score compared to decision tree.
2. Choosing the Best Model: Based on the evaluation metrics, the decision tree is preferred due to its balance between precision and recall and its higher ROC-AUC score. This model would be recommended to SyriaTel for predicting customer churn.

Final Recommendations to SyriaTel:

Develop Targeted Retention Strategies: Use the model to identify customers at high risk of churning based on key features like tenure, contract type, and monthly charges. Implement personalized retention strategies for these customers.

Enhance Customer Experience: Improve customer service quality and responsiveness, especially for those with frequent service interactions. Ensure issues are resolved quickly and effectively to prevent dissatisfaction from leading to churn.

Promote Long-Term Commitments: Encourage customers to move away from month-to-month contracts by offering incentives and making long-term contracts more flexible and appealing.

Optimize Pricing and Value Communication: Review pricing strategies, particularly for high-paying customers, and clearly communicate the value they receive for their payments.

Leverage Predictive Analytics: Continuously use predictive analytics to monitor customer behavior and identify emerging churn risks. Adjust retention strategies as new patterns emerge.